ChatGPT

ARCNet: Augmented Intelligence through Open AI Infrastructure

Autonomous Resource Controller Network (ARCNet) – a subsidiary of Arc Public Benefit Corp – is pioneering the world's first global **open-architecture, model-agnostic AI compute network**. This white paper presents ARCNet's vision and roadmap (2025–2028) and situates it within a scholarly historical context of *Augmented Intelligence*. We trace the lineage from early pioneers like Vannevar Bush and Douglas Engelbart, through Xerox PARC's breakthroughs, to modern open-source AI frameworks, grounding ARCNet's philosophy in a tradition that views technology as a means to **amplify human intellect rather than replace it** [1] [2]. We then examine the current AI landscape – namely the concentration of compute power in a few proprietary labs (OpenAI, Anthropic, Google DeepMind, etc.) – to underscore the urgency of ARCNet's open, public-benefit alternative. Finally, we detail how ARCNet's infrastructure enables **scalable, auditable, human-in-the-loop** AI systems and highlight the design, economic, and ethical advantages of ARCNet's approach over today's closed models.

Executive Summary

ARCNet is building a nationwide network of **10 hyperscale "AI Factory" data centers** (100 MW each) in partnership with NVIDIA [3]. This open platform will empower anyone – from researchers and startups to enterprises and domain experts – to build, train, and deploy large-scale AI models at **10× lower cost per token** than today's closed labs [4] [5]. Key features include **usage-based pricing with no artificial caps**, a model-agnostic architecture (users can run any model of their choice), and a sustainable energy strategy (mixing solar, battery storage, and advanced nuclear) for low, stable operating costs [6]. NVIDIA serves as an anchor investor providing cutting-edge hardware (e.g. next-generation 800 V DC data center design and "Rubin" class AI chips) while ARCNet retains independent control (Arc PBC owns 100% of ARCNet) to preserve its open mission [7]. ARCNet's public-benefit mandate allows it to prioritize **accessibility and affordability over profit**, operating with modest margins and passing efficiency gains to users [8]. In essence, ARCNet aims to set a new standard for open AI compute – analogous to how open-source software disrupted proprietary software – by **democratizing access** to advanced AI infrastructure [9]. This initiative directly challenges the status quo of AI compute being dominated by a few big tech companies, and instead treats **AI computation as a public utility**, available to all innovators without prohibitive costs or gatekeeping.

Historical Foundations of Augmented Intelligence

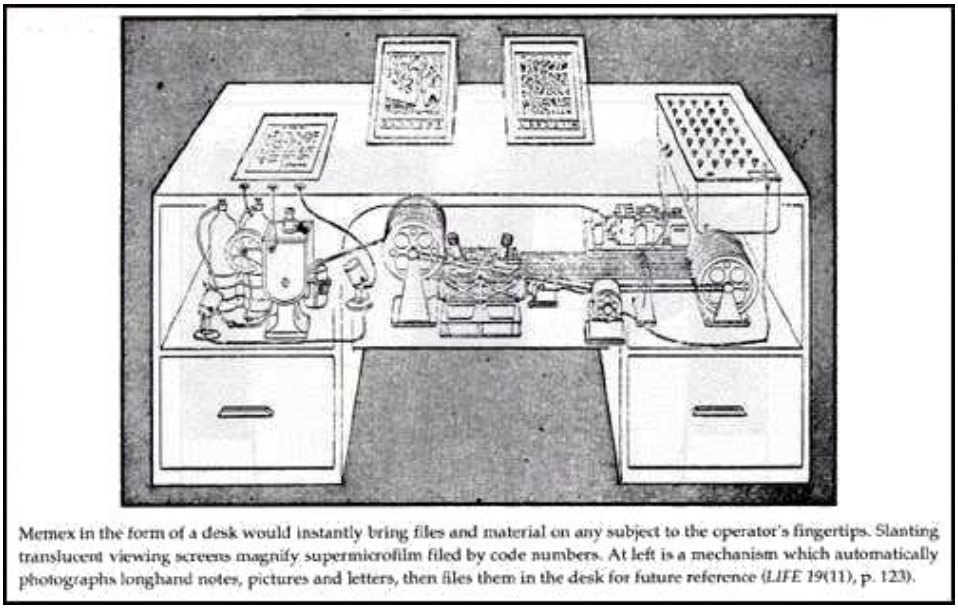

Memex in the form of a desk would instantly bring files and material on any subject to the operator's fingertips. Slanting translucent viewing screens magnify supermicrofilm filed by code numbers. At left is a mechanism which automatically photographs longhand notes, pictures and letters, then files them in the desk for future reference (*LIFE* 19(11), p. 123).

A 1945 illustration of Vannevar Bush's proposed Memex device, a "memory extender" envisioned as a personal knowledge base to augment human intellect (Life Magazine, 1945) [10] . In his seminal 1945 essay *"As We May Think,"* Vannevar **Bush** outlined a visionary plan for harnessing technology to amplify human intelligence and collective memory [11] . Alarmed that science was being directed toward war and destruction, Bush proposed the **Memex**, a hypothetical desk-like machine that would store vast amounts of information on microfilm and allow users to retrieve and link documents at will [12] [13] . This "collective memory machine" was intended to **transform the coming** information explosion **into a knowledge explosion** by making knowledge more accessible and navigable [14] . Bush's vision anticipated many aspects of modern computing – including hypertext linking, personal libraries, and associative search – decades before the digital computer era [13] . *As We May Think* deeply influenced the next generation of computing pioneers; notably, a young engineer named **Douglas Engelbart** read Bush's article and took it as a call to action [15] .

In the early 1960s, **Douglas Engelbart** began formulating a comprehensive framework for "augmenting human intellect." In a 1962 report titled *"Augmenting Human Intellect: A Conceptual Framework,"* Engelbart defined the goal succinctly: *"By augmenting human intellect we mean increasing the capability of [a person] to approach complex problem situations, to gain comprehension to suit [their] particular needs, and to derive solutions to problems."* [16] . Increased capability would entail **more rapid and better comprehension, speedier and better solutions, and the ability to address problems that were previously insoluble** [16] . Crucially, Engelbart emphasized this was *"not [about] isolated clever tricks"* or narrow efficiency hacks, but **a new way of life** in which human intuition ("hunches, cut-and-try, intangibles, and the human 'feel' for a situation") would **fruitfully coexist with powerful new concepts and high-powered electronic aids** [17] . He argued that as the world's problems grew more complex and urgent, only a concerted effort to augment human intellect would allow society to cope – an *"enlightened pursuit"* worthy of our best efforts [18] [19] .

Engelbart didn't stop at theory. At the Stanford Research Institute he assembled the Augmentation Research Center and, by 1968, unveiled a working prototype of his ideas. In what's now known as **"The Mother of All Demos"** (1968), Engelbart demonstrated the oN-Line System (NLS): a computer system

featuring a graphical interface, screen sharing, hyperlink text, the **first computer mouse**, and real-time collaboration. This live demo astonished the computing world. Engelbart's NLS showed how computers could be used *"to help knowledge workers perform better and faster,"* rather than just automate calculations [20]. As one historical account notes, Engelbart believed *"computers...should augment rather than replace the human intellect."* [1]. At the time, this interactive, human-centric approach was radically different from the prevailing focus on Artificial Intelligence (AI) aimed at **replacing human reasoning**. In fact, Engelbart's insistence on interactive, augmentative computing was initially met with skepticism – "90 percent of people thought he was a crackpot" and considered his interactive ideas "a waste of time" compared to the "really good stuff" being done in AI labs of the day [21]. Nevertheless, a few visionaries (like ARPA director Bob Taylor) championed Engelbart's work, and its impact spread. Engelbart's concepts **"fed out into Xerox PARC and then Apple to take over the world"**, as those skeptics were proven wrong when the personal computing revolution took off in the 1970s and '80s [21].

Xerox PARC (Palo Alto Research Center) in the 1970s was the crucible that turned Engelbart's augmentative vision into the foundations of modern personal computing. Many of PARC's researchers were directly inspired by NLS – indeed, the Xerox **Alto** computer (1973), the first GUI-based personal workstation, was conceived explicitly under the influence of Engelbart's work [22]. The Alto introduced the **graphical user interface (GUI)** with windows and icons, a mouse for direct manipulation, **WYSIWYG** text editing, and networked computing [23] [24]. Though the Alto was never sold commercially, it demonstrated how computing could amplify office productivity and creative work, aligning perfectly with the augmented intelligence ethos. Many of Alto's ideas migrated to **Apple** (after Steve Jobs famously visited PARC in 1979) and into mainstream products like the Apple Macintosh in 1984 [25]. The lineage from Bush to Engelbart to PARC is clear: **Bush's memex inspired Engelbart, Engelbart's NLS inspired PARC, and PARC's innovations laid the groundwork for human-centered computing for decades to come** [15] [22].

Fast-forward to the 21st century, and the principle of **augmenting human intellect** is alive in the rise of **open, collaborative AI tools**. The explosion of machine learning and AI in the 2010s initially saw many advances locked inside big corporate or academic labs, inaccessible to most. But the emergence of **open-source AI frameworks and communities** – such as **Hugging Face** – has dramatically democratized AI development. Hugging Face, often called the "GitHub of AI," began in 2016 and pivoted in 2018 to go "all in on open source." It built the **Transformers library** and the Hugging Face Hub, which in a few years *"cracked open AI's most powerful tools and handed them to the masses."* [26] With only a few lines of code, developers anywhere could use state-of-the-art language models that previously were confined to giant tech companies' labs [27] [26]. By **standardizing model sharing and providing a platform for collaboration**, Hugging Face tore down the walls that had made cutting-edge AI the province of "only elite labs and big tech players" [28]. In 2020, the Hugging Face Hub introduced *model cards*, versioning, and an easy-to-use repository for thousands of models – bringing a *"culture of openness and collaboration"* to AI that many had thought impossible in a competitive field [29]. Today the Hub hosts **over 300,000 models** and has fostered a vibrant community where **open science can outpace even the most secretive corporate labs**, proving that *"democratization can accelerate innovation"* and that **openness can scale faster than secrecy** [30]. This modern open-source movement in AI echoes Engelbart's ethos: it is fundamentally about **augmenting the many** rather than just empowering the few. It provides tools that allow a broad community of humans to apply their expertise, collaborate, and collectively improve AI systems – precisely the *networked human-centric* vision that underpins ARCNet.

The 2025 AI Landscape: Proprietary Giants vs. Open Collaboration

As of 2025, the development of frontier AI systems (like advanced large language models) has become heavily **concentrated in a small handful of industry labs**. A recent MIT Sloan analysis bluntly concludes: *"Tight control over access to specialized infrastructure and capabilities will likely result in a concentrated generative AI market controlled by a few key players."* [31] In practice, those key players are the well-known AI labs and their Big Tech sponsors: **OpenAI** (with Microsoft), **Anthropic** (with Google and Amazon backing), **Google DeepMind**, **Meta AI**, and a few others. These organizations command virtually exclusive access to the **massive compute resources, data, and funding** needed to train the largest models. For example, **OpenAI** alone reportedly *"spent billions of dollars on compute in 2024"* – roughly **$5 billion** on R&D (training) compute and **$2 billion** on inference compute – running on Microsoft's cloud supercomputers [32]. **Meta** has announced plans to build a dedicated AI supercomputer with **350,000 of NVIDIA's H100 GPUs by end of 2024**, an investment valued around **$10 billion** just in hardware [33]. Alphabet/Google and Microsoft, as cloud hyperscalers, are each investing on the order of **$80–90 billion** in AI data centers in 2025 alone [34] – staggering sums that no smaller player can hope to match. This arms race has led to what one commentator calls a *"multitrillion-dollar race"* between a few AI labs, with tech giants taking sides (Microsoft and NVIDIA backing OpenAI, Google and Amazon backing Anthropic, etc.) [35] [36]. For instance, Anthropic – founded by former OpenAI researchers – has secured over $4 billion from Amazon and over $3 billion from Google, and a recent cloud deal gives Anthropic access to **up to 1 million** of Google's AI-optimized TPU chips for its model training [37]. Amazon, meanwhile, is standing up one of the world's largest AI clusters (**Project Rainer** with ~**500,000** custom Trainium chips) dedicated to Anthropic's workloads [38]. These investments illustrate the **enormous concentration of AI compute in a few hands**.

Beyond compute concentration, there is also an **insularity of AI knowledge and control**. The most advanced models (like OpenAI's GPT-4, Google's Gemini, or Anthropic's Claude 2) are proprietary; outsiders can only access them through narrow APIs. As the MIT Sloan report notes, *"AI experts have quipped that the only thing open about OpenAI…is its name."* OpenAI has *"shrouded important technical details in secrecy, including model size, hardware, training compute, dataset construction, and training method."* [39] Without access to a model's weights or training data, **no external researcher or stakeholder can fully *replicate, audit, or improve* these systems** [40]. This black-box paradigm poses challenges for scientific transparency, accountability, and safety. Although the emergence of some open-source models (e.g. Meta's LLaMA, which was partially open-sourced in 2023) gave hope for a more "free-for-all" innovation environment, the reality is that **resources still dominate**. As Azoulay et al. (2024) point out, simply open-sourcing a model's code or weights is not enough to overcome the *"complementary assets"* that incumbents control [41] [42]. These include access to enormous **compute clusters**, massive **non-public datasets**, sophisticated evaluation pipelines, safety and compliance infrastructures, and the *network effects* of billions of users generating feedback data [43] [44]. The leading firms have a stranglehold on these assets by virtue of their head start and deep pockets, making it likely that **"new entrants…will need to rent access"** from those same firms if they wish to compete in building cutting-edge AI [45].

In short, the current trajectory of AI is **oligopolistic**: a few companies will control the core AI models and profit by gating their usage. An analogy can be drawn to the smartphone industry – where Apple and Google control the only two operating systems, and all developers must play by their rules [46] [47]. If AI goes the same way, we might see a world where **AI capabilities are ubiquitous but fundamentally *under the terms set by a small number of corporate gatekeepers*** [47]. This raises concerns not only about competition and pricing (monopoly rents on AI services), but also about **innovation stagnation, lack of transparency, and societal dependence** on entities that may not prioritize the public interest. As the Sloan

study warns, it's *"a lot of power in the hands of a very small number of companies."* [47] And indeed, **private industry now dominates AI research output**, eclipsing the role of academia [48]. This concentration has spurred calls for new approaches – whether through policy (antitrust and open competition) or through technology initiatives – to **"avert an oligopolistic future for the generative AI industry."** [49]

One promising sign is that **"going open"** can be a viable strategy, as demonstrated by Meta's partial openness with LLaMA. After LLaMA's weights leaked in 2023, Meta chose to officially open-source a version (Llama 2) in 2024, with CEO Mark Zuckerberg arguing that an open approach is *"the path forward"* for AI, benefiting academic and independent researchers [50]. This was a strategic move to gain community goodwill and adoption – something ARCNet takes to heart. However, relying on the goodwill of a single tech giant is risky; even Meta could change course if openness no longer suits its interests [51]. The broader lesson is that **genuine openness in AI needs robust, independent infrastructure** and organizations **explicitly committed to the public interest**. This is precisely the gap ARCNet intends to fill. ARCNet enters this landscape as an **unprecedented public-benefit AI utility** – one that seeks to decentralize and democratize the compute power and capabilities that have so far been the province of a few firms.

ARCNet's Vision: An Open, Human-Centric AI Network

ARCNet is founded on the principle of Augmented Intelligence: networking humans with AI to amplify collective expertise, not to replace human judgment or creativity. In spirit, ARCNet is a direct inheritor of the Engelbartian vision – combining powerful "electronic aids" with human intellect to tackle complex problems [17]. Just as Engelbart targeted the "professional problems of diplomats, executives, scientists, attorneys, designers" with his augmentation framework [52], ARCNet's mission is to empower experts across **science, medicine, law, education**, and beyond with advanced AI tools that they **control and direct**. *"Automation shouldn't mean abdication,"* as modern AI ethicists say – the goal is AI **with** humans in the loop, where *"the power of AI lies in scale"* and *"the wisdom of humans lies in context."* [53] ARCNet's approach ensures that **AI makes suggestions, not final decisions; humans can always override or refine the outputs; and continuous learning is a two-way street** (AI helps humans discover insights faster, while humans guide AI with feedback) [54]. By keeping people at the center of the design, ARCNet is aligning with the fundamental truth that *"AI isn't here to replace us; it's here to work with us"* – technology *"that enhances people, rather than marginalizing them."* [2]

Concretely, ARCNet is building the **infrastructure and governance** to make this human-centric model possible at scale. Unlike proprietary AI platforms, ARCNet's **open architecture** means users are not locked into specific models or software ecosystems. Developers and researchers will have the freedom to deploy **any model** (open-source or their own) on ARCNet's compute network, with full access to the model's workings. This model-agnostic design, coupled with transparent operations, makes **auditing and reproducibility** feasible. For example, an academic team could train a model on ARCNet and share not just the final model but the entire training log and configuration – anyone could inspect how it was trained, identify biases, or replicate the experiment. This level of **auditability and transparency** is a stark contrast to closed labs where such details are proprietary secrets [55].

Moreover, ARCNet's **usage-based pricing** and public-benefit orientation remove the financial and artificial barriers that currently limit who can leverage AI. There will be *"no hard limits or waitlists"* for access – users pay only for the compute they actually use, at reasonable rates with modest margins [56]. A **student or startup** can thus access thousands of GPUs on ARCNet for a short time to run an experiment, without needing venture capital or special partnerships (whereas today, large-scale experiments are often only

possible if you are inside a big tech firm or have cloud credits from one). This is analogous to cloud computing's pay-as-you-go model but specialized for AI and without the monopolistic pricing power of a few hyperscalers. ARCNet essentially treats **AI computation as a public utility** – much like electricity or internet bandwidth – that should be available to all at fair prices [57] [8]. By 2028, with 10 AI factories online, ARCNet projects an impressive nationwide capacity that can support **the broadest community of AI developers and domain experts** in the world.

Infrastructure for Scalable, Auditable, Human-in-the-Loop AI

ARCNet's **technology infrastructure** is deliberately designed for **scale, efficiency, and openness**. Each of the 10 planned data centers (the "AI factories") is a **100 MW facility**, built in partnership with NVIDIA's engineering team [3]. These centers use NVIDIA's state-of-the-art reference designs circa 2025–2027, such as a full **800 V DC power distribution architecture** [58] [59]. By eliminating legacy AC-DC power conversions, the 800 V DC design improves power delivery efficiency by ~150% and allows an unprecedented density of GPUs per rack [60] [58]. In practical terms, this means ARCNet can pack more computing power into each facility at lower operating cost, directly translating to **better performance-per-dollar for users** [61] [60]. Each site will host tens of thousands of cutting-edge GPUs or AI accelerators, networked with high-bandwidth interconnects to function as one giant distributed supercomputer. The network of sites will eventually interconnect, creating a **continental-scale AI compute grid**. This massive scale is what enables ARCNet to offer **nation-level AI capability** on an open-access basis.

Importantly, ARCNet couples scale with **reliability and auditability**. All systems are being built with **robust monitoring, logging, and controls** such that every job run on ARCNet can produce an **"audit trail."** From data provenance (tracking which datasets were used to train a model) to model evaluation results, the platform is being optimized for **traceability**. This focus is driven by both practical and ethical needs: in sectors like healthcare or law, practitioners will require **explainability and records** of how an AI-assisted recommendation was generated. ARCNet's open approach means that, unlike a black-box API, the internal workings of models and computations are not off-limits to the user. An ARCNet user can choose or develop a model that is fully transparent (e.g. open-source weights), and run it in an environment where every parameter change or output can be logged. This makes **third-party audits and compliance checks** far easier. For instance, a hospital using ARCNet to run an AI diagnosis tool can retain logs needed for regulatory compliance, and even invite external auditors to inspect the model for biases or errors – a process that would be impossible if the model were a proprietary service. By design, ARCNet aligns with emerging best practices in **AI ethics and governance**, which call for *open data, AI audits, and explainability* especially in high-stakes domains [62] [63].

Another cornerstone of ARCNet's infrastructure is support for **human-in-the-loop workflows** at scale. This means the platform will natively support modes where AI systems and humans interact frequently and intentionally. For example, ARCNet can enable **iterative model refinement**, where a domain expert runs a model, examines its output, and then adjusts the model or provides feedback before the next run. Many current AI deployment pipelines are one-shot (train -> deploy -> full automation); by contrast, ARCNet's tooling encourages **continuous supervision and improvement**. In a sense, ARCNet treats AI models not as autonomous oracles but as **collaborative assistants** that remain connected to human operators. The system could allow, say, a scientist to "peek" at a model's intermediate reasoning steps or uncertainties and intervene as needed. It could also facilitate **crowd-sourced expertise**: consider a legal AI system on ARCNet that flags a set of ambiguous cases and routes them to a network of volunteer legal experts for review – their feedback then updates the model. Because ARCNet is open and not tied to a single company's product

timeline, it can explore these *human-in-loop at scale* paradigms in a way that closed services (which often seek full automation for profitability) might not.

Example Use Cases Enabled by ARCNet

- **Scientific Research:** A multi-institution research team in climate science is investigating complex Earth system models. Using ARCNet, they train a large AI model to emulate certain climate simulations. The **entire process is auditable** – every training epoch's results and parameters are logged. Scientists remain in the loop by periodically reviewing the model's predictions against known benchmarks; if the model drifts from physical plausibility, they intervene with new training data or adjusted parameters. **Collectively, dozens of scientists around the world collaborate on ARCNet** – sharing compute resources, data, and model improvements in an open repository. This setup amplifies their research capabilities, allowing them to explore climate scenarios 10× faster than before, while ensuring the process is transparent and scientifically valid.

- **Healthcare (Medicine):** A medical center uses ARCNet to develop a clinical AI assistant for radiology. Rather than adopting a pre-packaged black-box model, their doctors and data scientists fine-tune an open-source medical imaging model on ARCNet with the hospital's own anonymized scans. The system is **human-in-the-loop**: the AI suggests diagnostic annotations on imaging studies, but a radiologist reviews each suggestion. Over time, the radiologists give feedback (confirming or correcting the AI's suggestions), and those labels are fed back on ARCNet to continually improve the model. All actions are **auditable**, which is crucial for compliance with health regulations – the hospital can trace exactly which images and doctor feedback led to each model update. The result is a highly accurate AI that *augments* the radiologists' workflow (speeding up measurements, highlighting anomalies) while the physicians maintain full control and final judgment on each case. Patient outcomes improve, and the doctors trust the system because it was built *by them, for them* on an open platform, with no opaque "AI magic."

- **Law and Justice:** Public defense attorneys form a consortium to create an AI legal research assistant on ARCNet, aimed at helping with case preparation. They train language models on a large corpus of public legal documents (statutes, court opinions, briefs) – all of which can be hosted in ARCNet's open data libraries. The **AI model remains open and scrutable**: attorneys can see which precedents the model is citing and how it derives its suggestions. During use, the AI might draft a memorandum or suggest relevant cases, but **the lawyer is always in the loop to edit and approve**. If the model makes an incorrect citation or odd suggestion, the attorney corrects it, and this feedback is recorded to avoid similar errors in the future. Because the system is open, independent evaluators (say, a civil liberties group) could audit it for biases (e.g., ensure it's not systematically disadvantaging certain groups in its recommendations) – something that would be impossible with a closed AI service. The **collective knowledge of many legal experts** effectively gets distilled into the model via ARCNet, creating a resource that improves access to justice (small public defender offices now have research support that rivals large law firms' capabilities) without replacing the essential human judgment of lawyers.

- **Education:** A group of teachers and educational researchers collaborate on ARCNet to build AI tutoring systems for K-12 students. Using open-source language models, they customize them with pedagogical strategies and local curriculum content. The **tutors run on ARCNet** and are deployed in classrooms in a *human-supervised* manner: an AI tutor might give a student practice problems and

feedback, but teachers oversee the process through ARCNet's dashboard, which shows in real time how the AI is interacting and where a student might be struggling. Teachers can **intervene at any moment**, for example by adjusting the difficulty of questions the AI presents or by providing a human explanation when needed. All interactions are logged (with student privacy protections) so that educators can later analyze and audit the AI's behavior: Was it fair? Did it stick to approved materials? This data also helps improve the tutoring model continuously. The outcome is a scalable personalized learning tool that **amplifies teachers' reach** – each student gets more tailored practice – yet the system is aligned with human educational values because teachers remain deeply involved. The open nature of ARCNet means the **curriculum designers and child psychologists** can also inspect and tweak the tutor's algorithms, ensuring they meet educational goals and ethical standards.

Through these examples (and many more use cases across domains), ARCNet enables *"a way of life in an integrated domain"* where **human expertise, intuition, and ethics coexist with powerful AI tools** [17]. The networked structure of ARCNet means **collective human knowledge can be brought to bear like never before** – e.g. scientists in different countries pooling their know-how on one platform, or niche expert communities (like rare disease specialists) developing their own AI models and sharing them openly. This human-centric philosophy is baked into ARCNet's design, fulfilling Engelbart's ideal of boosting society's problem-solving capacity by *connecting humans and computers in synergistic loops*.

Design, Economic and Ethical Benefits of ARCNet's Open Model

Design Benefits (Open & Modular): ARCNet's open architecture stands in contrast to the siloed, product-specific designs of proprietary AI services. By being **model-agnostic and open-architecture**, ARCNet can rapidly incorporate the best ideas from anywhere – whether it's a new open-source model from the community or a specialized hardware accelerator. This flexibility accelerates innovation: users aren't waiting on a vendor's roadmap; they can bring their own advancements to ARCNet. The design also emphasizes **standards and interoperability**, making it easier to audit and integrate. For example, ARCNet might adopt standard formats for model weights, training data lineage, and evaluation metrics, so that contributions from different groups remain compatible and reviewable. The influence of open-source software design is deliberate: just as open protocols and OSes led to a flourishing of software innovation, an open AI compute design can unleash creativity and trust. Jensen Huang, NVIDIA's CEO, has described next-gen AI data centers as "**AI factories**" – ARCNet embodies this but as *open* AI factories, where the schematics aren't secret [64]. This openness is **analogous to how Linux and open-source platforms disrupted proprietary systems** – by inviting collaboration and inspection, ARCNet can actually **achieve greater reliability and security** (issues are identified and fixed by the community) compared to closed systems that rely on security-through-obscurity. In short, ARCNet's design prioritizes **user empowerment, transparency, and adaptability**, ensuring that the infrastructure serves the users' goals, not the other way around.

Economic Benefits (Access & Efficiency): ARCNet's economic model is built around **democratization of access** and cost efficiency. Because it is structured as a Public Benefit Corporation and not a profit-maximizing entity, ARCNet can operate on much thinner margins than typical cloud providers or AI companies [8]. This means savings are passed directly to users – the target of **10× lower cost-per-token** for AI training/inference is transformative [5]. Such cost reduction lowers the barrier to entry for cash-strapped researchers, startups, or public institutions like universities and hospitals. In effect, it levels the playing field: innovative ideas can be tested and scaled without needing access to a corporate war chest. The usage-based pricing with **no artificial caps or tiering** also prevents the common scenario of large

incumbents throttling usage or charging premium fees for higher-end capabilities [56]. Economically, ARCNet treats AI compute as a commodity utility – abundant and cheap – rather than a scarce luxury. This could unleash a wave of AI adoption and experimentation in sectors that so far have been left behind (e.g., small businesses or developing countries' institutions that couldn't afford big cloud bills). Moreover, ARCNet's emphasis on **efficiency** (e.g. advanced power and cooling, high utilization of hardware, strategic siting of data centers for cheap renewable energy) means the absolute cost of providing compute is lower. By partnering with NVIDIA and using the latest hardware at scale, ARCNet achieves economies of scale comparable to the tech giants, but with the **critical difference that the resulting capacity is available to all**. The broader economic benefit is an **increase in AI-driven innovation and productivity outside the Big Tech sphere**, potentially leading to more diverse AI applications, new startups, and community-driven research breakthroughs. This diffusion of innovation can have positive spillovers on the economy, analogous to how open internet infrastructure sparked the dot-com boom by being available to all entrepreneurs.

Ethical and Societal Benefits (Accountability & Equity): ARCNet's open, public-benefit model inherently addresses many ethical concerns prevalent in AI today. First, **accountability and transparency**: an AI system built and deployed on ARCNet can be audited by design – stakeholders (be it users, regulators, or affected communities) can scrutinize how it works [55]. This is vital for building trust in AI. For instance, bias in AI outputs can be identified and corrected in the open, rather than remaining a mystery. If an AI model is used in government or healthcare decisions, the open platform ensures it can be independently evaluated for fairness or accuracy. Second, **human agency** is preserved. ARCNet explicitly rejects the "fully autonomous AI" narrative in favor of *human-in-the-loop* systems; this aligns with emerging ethical frameworks that insist on human final authority especially in sensitive applications (the EU's draft AI Act, for example, emphasizes human oversight in high-risk AI). By keeping humans involved and giving them control, ARCNet mitigates the risk of blindly algorithmic decisions and the erosion of human skills. Third, ARCNet's **public-benefit governance** is an ethical innovation in itself. It is **accountable not to shareholders but to its mission** – which is to serve the public good by widening AI access. This reduces conflicts of interest that plague for-profit AI providers (e.g., the incentive to hoard data, or to prioritize engagement over well-being). ARCNet can make decisions like open-sourcing a model or throttling a dangerous capability based on societal benefit, not profit. Additionally, by diversifying the AI ecosystem, ARCNet helps counteract the **centralization of power**. No single company or nation should unilaterally decide the "rules" of advanced AI; an open network dilutes power and invites **global participation and oversight**. Finally, there are potential **environmental and social benefits**: ARCNet's commitment to sustainable energy (mixing solar, batteries, and even small modular reactors) for its compute means the carbon footprint per AI operation is minimized [65]. This is significant as AI workloads grow in scale – ARCNet can set a benchmark for green AI compute. Socially, if ARCNet succeeds, the benefits of AI (in healthcare outcomes, education, scientific discovery, economic growth) will be more widely and equitably distributed, rather than concentrated in a few tech hubs. In summary, ARCNet's approach fosters an **AI ecosystem that is more accountable, inclusive, and aligned with human values** compared to the status quo of proprietary AI development.

Conclusion

In establishing ARCNet, Arc Public Benefit Corp is operationalizing a bold idea: that **the future of AI should belong to everyone**, and that the proper role of AI is as an amplifier of human ingenuity, not a replacement. This philosophy harks back to the earliest dreams of computing pioneers – Vannevar Bush's collective knowledge machine, Engelbart's human-computer symbiosis, and the communal spirit of early

internet and open-source communities. Engelbart wrote in 1962 of *"a way of life in an integrated domain"* where humans, with their intuition and creativity, work in partnership with powerful tools to solve the problems of their time [17]. ARCNet is a 21st-century realization of that vision: a nationwide infrastructure explicitly designed to **augment human intellect at the collective level**, by linking individuals and institutions into a shared network of knowledge and capabilities.

This white paper has outlined how ARCNet's approach – historically grounded, technologically ambitious, and ethically committed – provides a compelling alternative to the AI oligopoly that is emerging. By **2028**, when ARCNet's ten AI factories are fully operational across the U.S., we foresee a thriving ecosystem of innovation radiating from them: startups building novel applications without needing permission or huge capital, researchers tackling grand challenges (from climate modeling to curing diseases) with unprecedented compute at their fingertips, and professionals in fields like medicine, law, and education leveraging AI in a way that **enhances their expertise and effectiveness** rather than deskilling them. Crucially, because ARCNet is open and auditable, society at large can **trust** and shape these developments. Regulators, academics, and citizen groups will have access to the "innards" of the important AI systems affecting their lives, enabling informed dialogue and oversight that would be impossible behind corporate silos.

In the end, ARCNet's significance may be as much *cultural* as it is technical or economic. It asserts that there is an alternate path for AI – one rooted in **cooperation, transparency, and public good** – which stands in contrast to the competitive, secretive, and profit-driven trajectory that has dominated recent years. By proving the viability of a large-scale open AI utility, ARCNet could catalyze more initiatives around the world (imagine "ARCNet Europe" or networks in Asia/Africa dedicated to local needs). The network effect here is not just silicon-based, but human: **networking humans via AI** to create a super-intelligence that is *our* intelligence, amplified [16] [2]. This is a future where AI is not an alien force or proprietary product, but an **integral, transparent extension of the human collective**, helping us to **think faster, comprehend more, and solve what once seemed unsolvable** [16]. It is an aspirational vision – but one that ARCNet's plan, grounded in both historical insight and cutting-edge engineering, is poised to deliver.

References:

1. Engelbart, D.C. (1962). *Augmenting Human Intellect: A Conceptual Framework*. SRI Summary Report AFOSR-3223. (Excerpts: [16] [17] [52])

2. Bush, V. (1945). *As We May Think*. *The Atlantic*, July 1945. (Summary: Bush's memex concept for collective memory [11] [13])

3. Adam Fisher, *Wired* (2018). "How Doug Engelbart Pulled off the Mother of All Demos." (Engelbart's philosophy that computers should *augment* not replace human intellect [1]; historical perspective on NLS influence [21])

4. **MIT Sloan Management Review** (Brian Eastwood, 2024). "Openness, control, and competition in the generative AI marketplace." (Analysis of AI market concentration and need for open infrastructure [31] [44] [47])

5. You, J. (Epoch, 2025). "Most of OpenAI's 2024 compute went to experiments." (OpenAI's estimated $7B compute spend in 2024) [32]

6. Business Insider (Bergman, 2025). “Google in talks to deepen investment in Anthropic…” (Details on Anthropic’s funding, 1M TPU deal, Amazon cluster) [37] [38]

7. Philip J. Basile, *Dev.to* (2023). “The Rise of Hugging Face: How One Company Democratized AI…” (Describes how Hugging Face opened AI tools to the masses [26] [29])

8. O. Idowu, *Medium* (2024). “Embedding Ethics into AI Collaboration – The Future of Work.” (Discusses human-in-the-loop principles and trust design [54] [2])

9. **Autonomous Resource Controller Network (ARCNet)** – Arc PBC internal White Paper/Business Plan (2025). (Details on ARCNet vision, NVIDIA partnership, open model vs closed labs, etc. [8] [56])

10. Wikipedia – *Xerox Alto* (accessed 2025). (Historical note: Alto inspired by Engelbart’s work) [22]

11. Wikipedia – *As We May Think* (accessed 2025). (Historical note: Bush’s influence on Engelbart and future computing) [15]

---

1 20 21 How Doug Engelbart Pulled off the Mother of All Demos | WIRED
https://www.wired.com/story/how-doug-engelbart-pulled-off-the-mother-of-all-demos/

2 53 54 63 Embedding Ethics into AI Collaboration = The Future of Work | by Oyetola florence Idowu | Medium
https://medium.com/@oyetola13/embedding-ethics-into-ai-collaboration-the-future-of-work-554bd3b5dc3b

3 4 5 6 7 8 9 56 57 58 59 60 61 64 65 Autonomous Resource Controller Network (ARCNet).pdf
file://file_00000000f1bc71f59b8ffbd1294c2eec

10 Picture Gallery: Vannevar Bush's Memex, Drawing
https://www.kerryr.net/pioneers/gallery/ns_bush8.htm

11 12 13 14 15 As We May Think - Wikipedia
https://en.wikipedia.org/wiki/As_We_May_Think

16 17 18 19 52 Doug_Engelbart-AugmentingHumanIntellect.pdf
file://file_000000003d4471f58e0bc68a28ace4f8

22 23 24 25 Xerox Alto - Wikipedia
https://en.wikipedia.org/wiki/Xerox_Alto

26 27 28 29 30 The Rise of Hugging Face: How One Company Democratized AI and Transformed Machine Learning - DEV Community
https://dev.to/philipjohnbasile/the-rise-of-hugging-face-how-one-company-democratized-ai-and-transformed-machine-learning-2lmp

31 33 39 40 41 42 43 44 45 46 47 48 49 50 51 55 Openness, control, and competition in the generative AI marketplace | MIT Sloan
https://mitsloan.mit.edu/ideas-made-to-matter/openness-control-and-competition-generative-ai-marketplace

32 Most of OpenAI’s 2024 compute went to experiments | Epoch AI
https://epoch.ai/data-insights/openai-compute-spend

34 How Many AI Companies are There? [2026] | StartUs Insights
https://www.startus-insights.com/innovators-guide/how-many-ai-companies-are-there/

35 36 37 38 Google Is in Early Talks to Deepen Its Investment in Anthropic - Business Insider
https://www.businessinsider.com/google-deepen-investment-in-ai-anthropic-2025-11

62 Trust, Regulation, and Human-in-the-Loop AI
https://cacm.acm.org/research/trust-regulation-and-human-in-the-loop-ai/